

TECH
VAULT



JOIN ALGORITHMS

- NESTED LOOP
- INDEX NESTED LOOP
- MERGE
- HASH

AMR ELHELW

All Here is about Inner join

Customers

cid	cname	email	Region
c1	Alice	alice@example.com	West
c2	Bob	bob@example.com	East
c3	Charlie	charlie@example.com	North
c4	Danny	danny@example.com	East

Orders

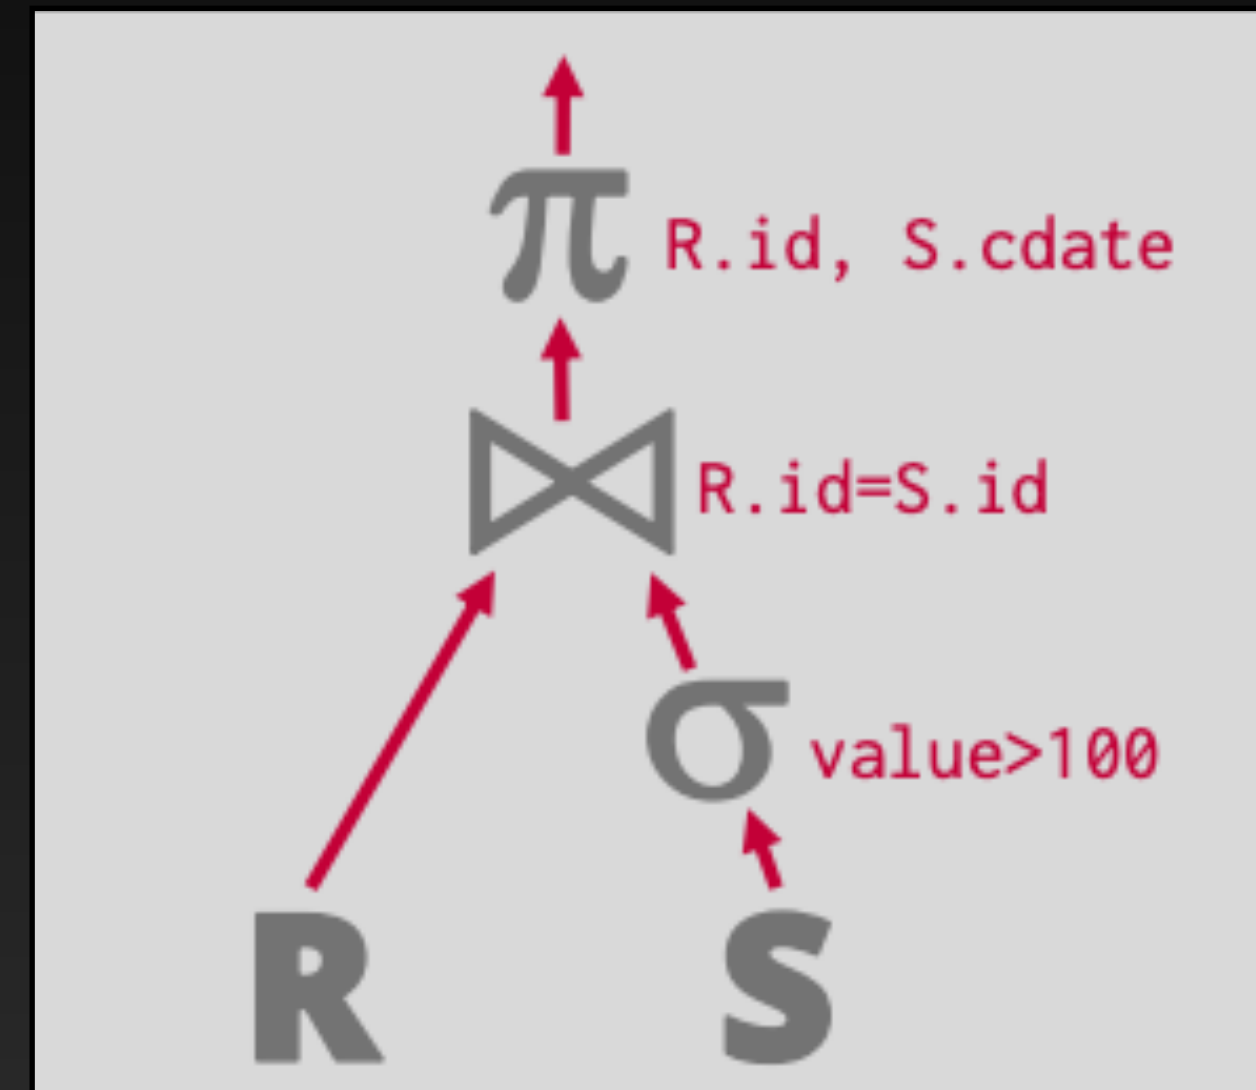
oid	odate	cid	amount
101	2022-04-10	c2	\$150
102	2023-05-07	c1	\$20
103	2021-10-04	c4	\$37
104	2021-10-04	c1	\$126
105	2023-06-20	c1	\$30
106	2022-07-25	c4	\$74

“Return all customer names and order amounts from the East region”

Order of joins matters in the cost.

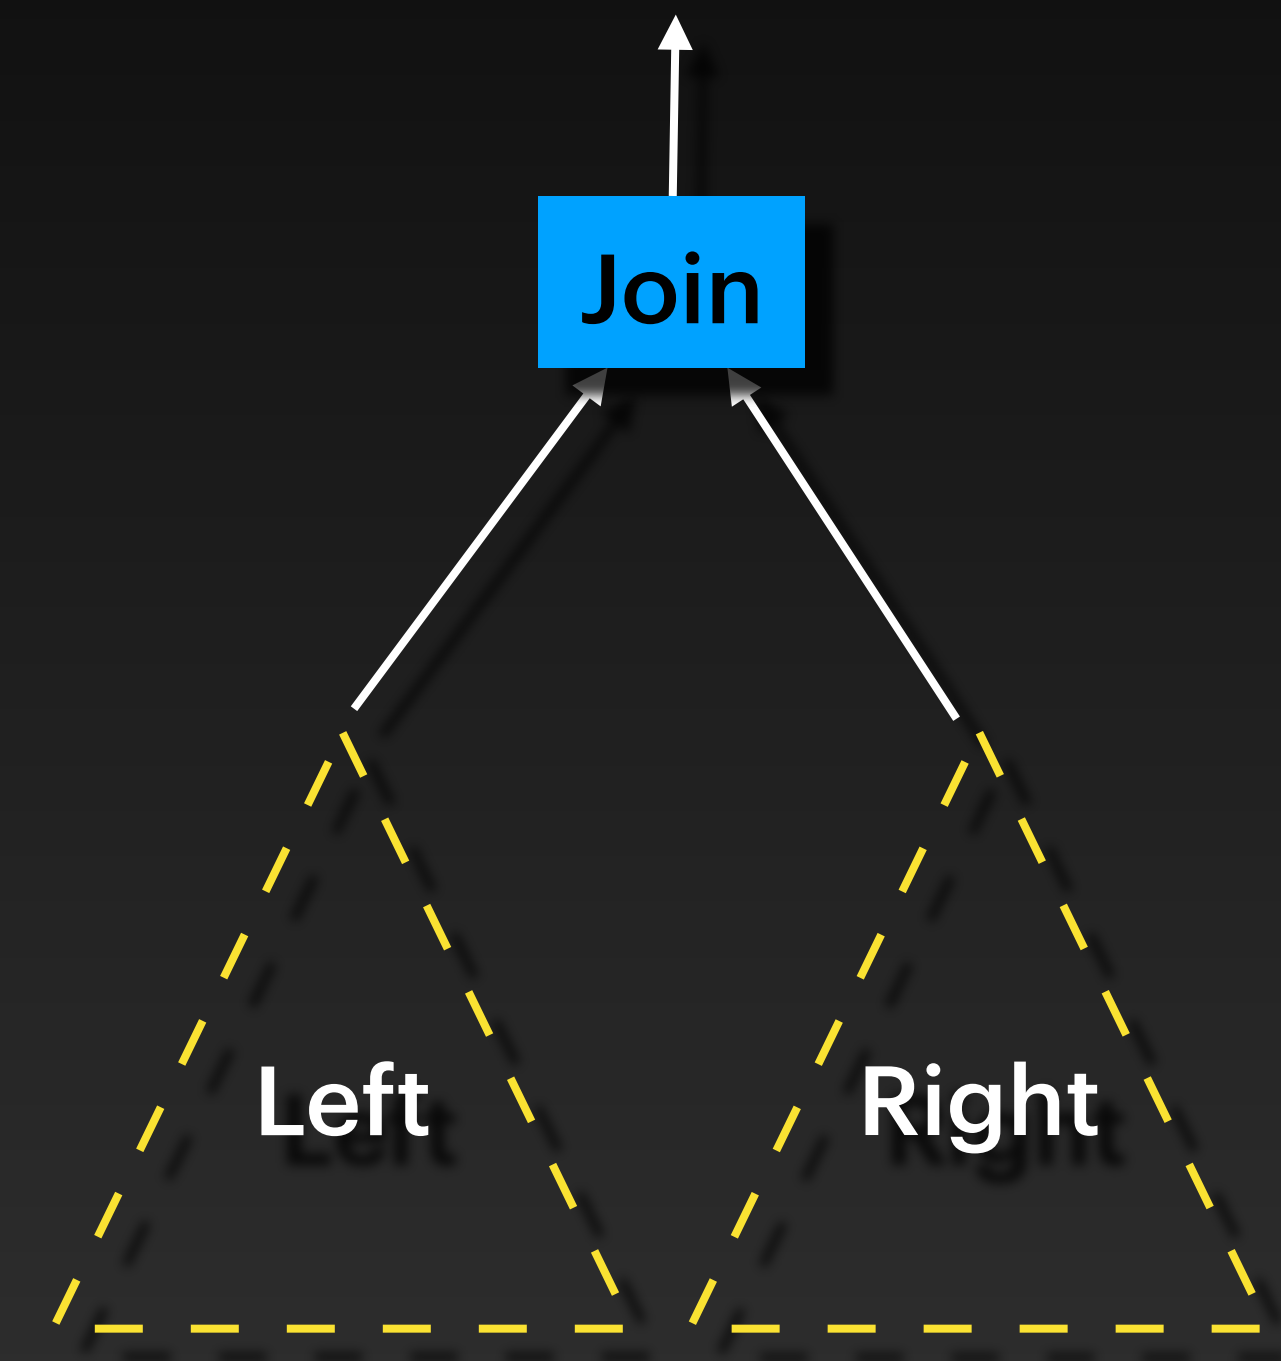
Join Query

```
SELECT R.id, S.cdate  
FROM R JOIN S  
ON R.id = S.id  
WHERE S.value > 100
```



Join Plan


- Each input can be a subplan
 - Left (outer)
 - Right (inner)
- Each join algorithm has a different cost
- Cost can be estimated in terms of number of rows fetched from each side



Nested Loop Join

```
SELECT *  
FROM R JOIN S  
ON R.id = S.id
```

```
foreach tuple r in R:  
    foreach tuple s in S:  
        if (condition): output row
```



R

id	name
3	Alice
2	Bob
1	Charlie
4	Danny

S

id	Value
2	1000
3	2000
4	1300
2	5000
4	2500

Output

R.id	R.name	S.id	S.value
3	Alice	3	2000
2	Bob	2	1000
2	Bob	2	5000
4	Danny	4	1300
4	Danny	4	2500

Generally:

- * nested loop join is the worst one because of nested.
- * is usually used in inequy joins.


Nested Loop Join

```
SELECT *  
FROM R JOIN S  
ON R.id = S.id
```

N_R : #rows in R


N_S : #rows in S

General Rule:
the lowest number of rows table
is better to be the outer table.
to minimize the left part.



$$\text{Cost} = N_R + (N_R * N_S)$$

if $N_R < N_S$ So 1st one is better
else second one is better
because the nested is fixed in both equations



$$\text{Cost} = N_S + (N_S * N_R)$$

Better to use the smaller input on the left side!

Nested Loop Join

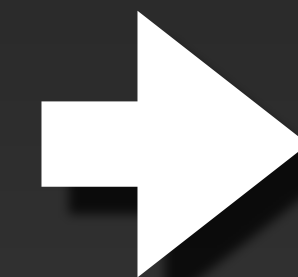
```
SELECT *  
FROM R JOIN S  
ON R.id = S.id
```

```
foreach tuple r in R:  
    foreach tuple s in S:  
        if (condition): output row
```

Search for tuples in *S*
with *id = r.id*

What if there was already an index on *S.id*?


```
foreach tuple r in R:  
    SS = Search Index for r.id  
    foreach tuple s in SS:  
        output row
```



Index Nested
Loop Join

Index Nested Loop Join

```
SELECT *  
FROM R JOIN S  
ON R.id = S.id
```



$$\text{Cost} = N_{\text{outer}} + (N_{\text{outer}} * \text{Cost_index_inner})$$

R

→

id	name
3	Alice
2	Bob
1	Charlie
4	Danny

S

id	Value
2	1000
3	2000
4	1300
2	5000
4	2500

Output

R.id	R.name	S.id	S.value
3	Alice	3	2000
2	Bob	2	1000
2	Bob	2	5000
4	Danny	4	1300
4	Danny	4	2500

Cost = N_outer + N_inner + Sort (if needed)

Merge Join

- * works only with equi joins
- * needs tables to be sorted on join key.
other wise it sorts them which is more c.ost.

```
SELECT *  
FROM R JOIN S  
ON R.id = S.id
```

perfect: if tables are sorted based on the join key.
how they can be sorted? if they comes from index scan, sorted before in a merge join, cached in a sort operation before

NOTE: the output is sorted which means it can be used later in another merge join.
and it is also useful if you need to order by that key.

R

id	name
3	Alice
2	Bob
3	Alice
4	Danny

S

id	Value
2	1000
3	2000
4	1300
4	5300
4	2500

Output

R.id	R.name	S.id	S.value
2	Bob	2	1000
2	Bob	2	5000
3	Alice	3	2000
4	Danny	4	1300
4	Danny	4	2500

HERE:

it sorts the 2 tables on id the:
it uses 2 pointers technique to move pointers
if both are equal. other wise move the less pointer and so on.

```
SELECT *  
FROM R JOIN S  
ON R.id = S.id
```

If there was already an index on S.id

```
foreach tuple r in R:  
    SS = Search Index for r.id  
    foreach tuple s in SS:  
        output row
```

If there was no index S.id

Build an “index” on the fly?

Focus on equi-joins —> Hash index

```
Build Hash table on S.id  
foreach tuple r in R:  
    SS = Search Hash table for r.id  
    foreach tuple s in SS:  
        output row
```

STEPS:

1. choose one table to build hash table on it.
2. iterate over the second table and lookup in the hash table.

Index Nested Loop Join

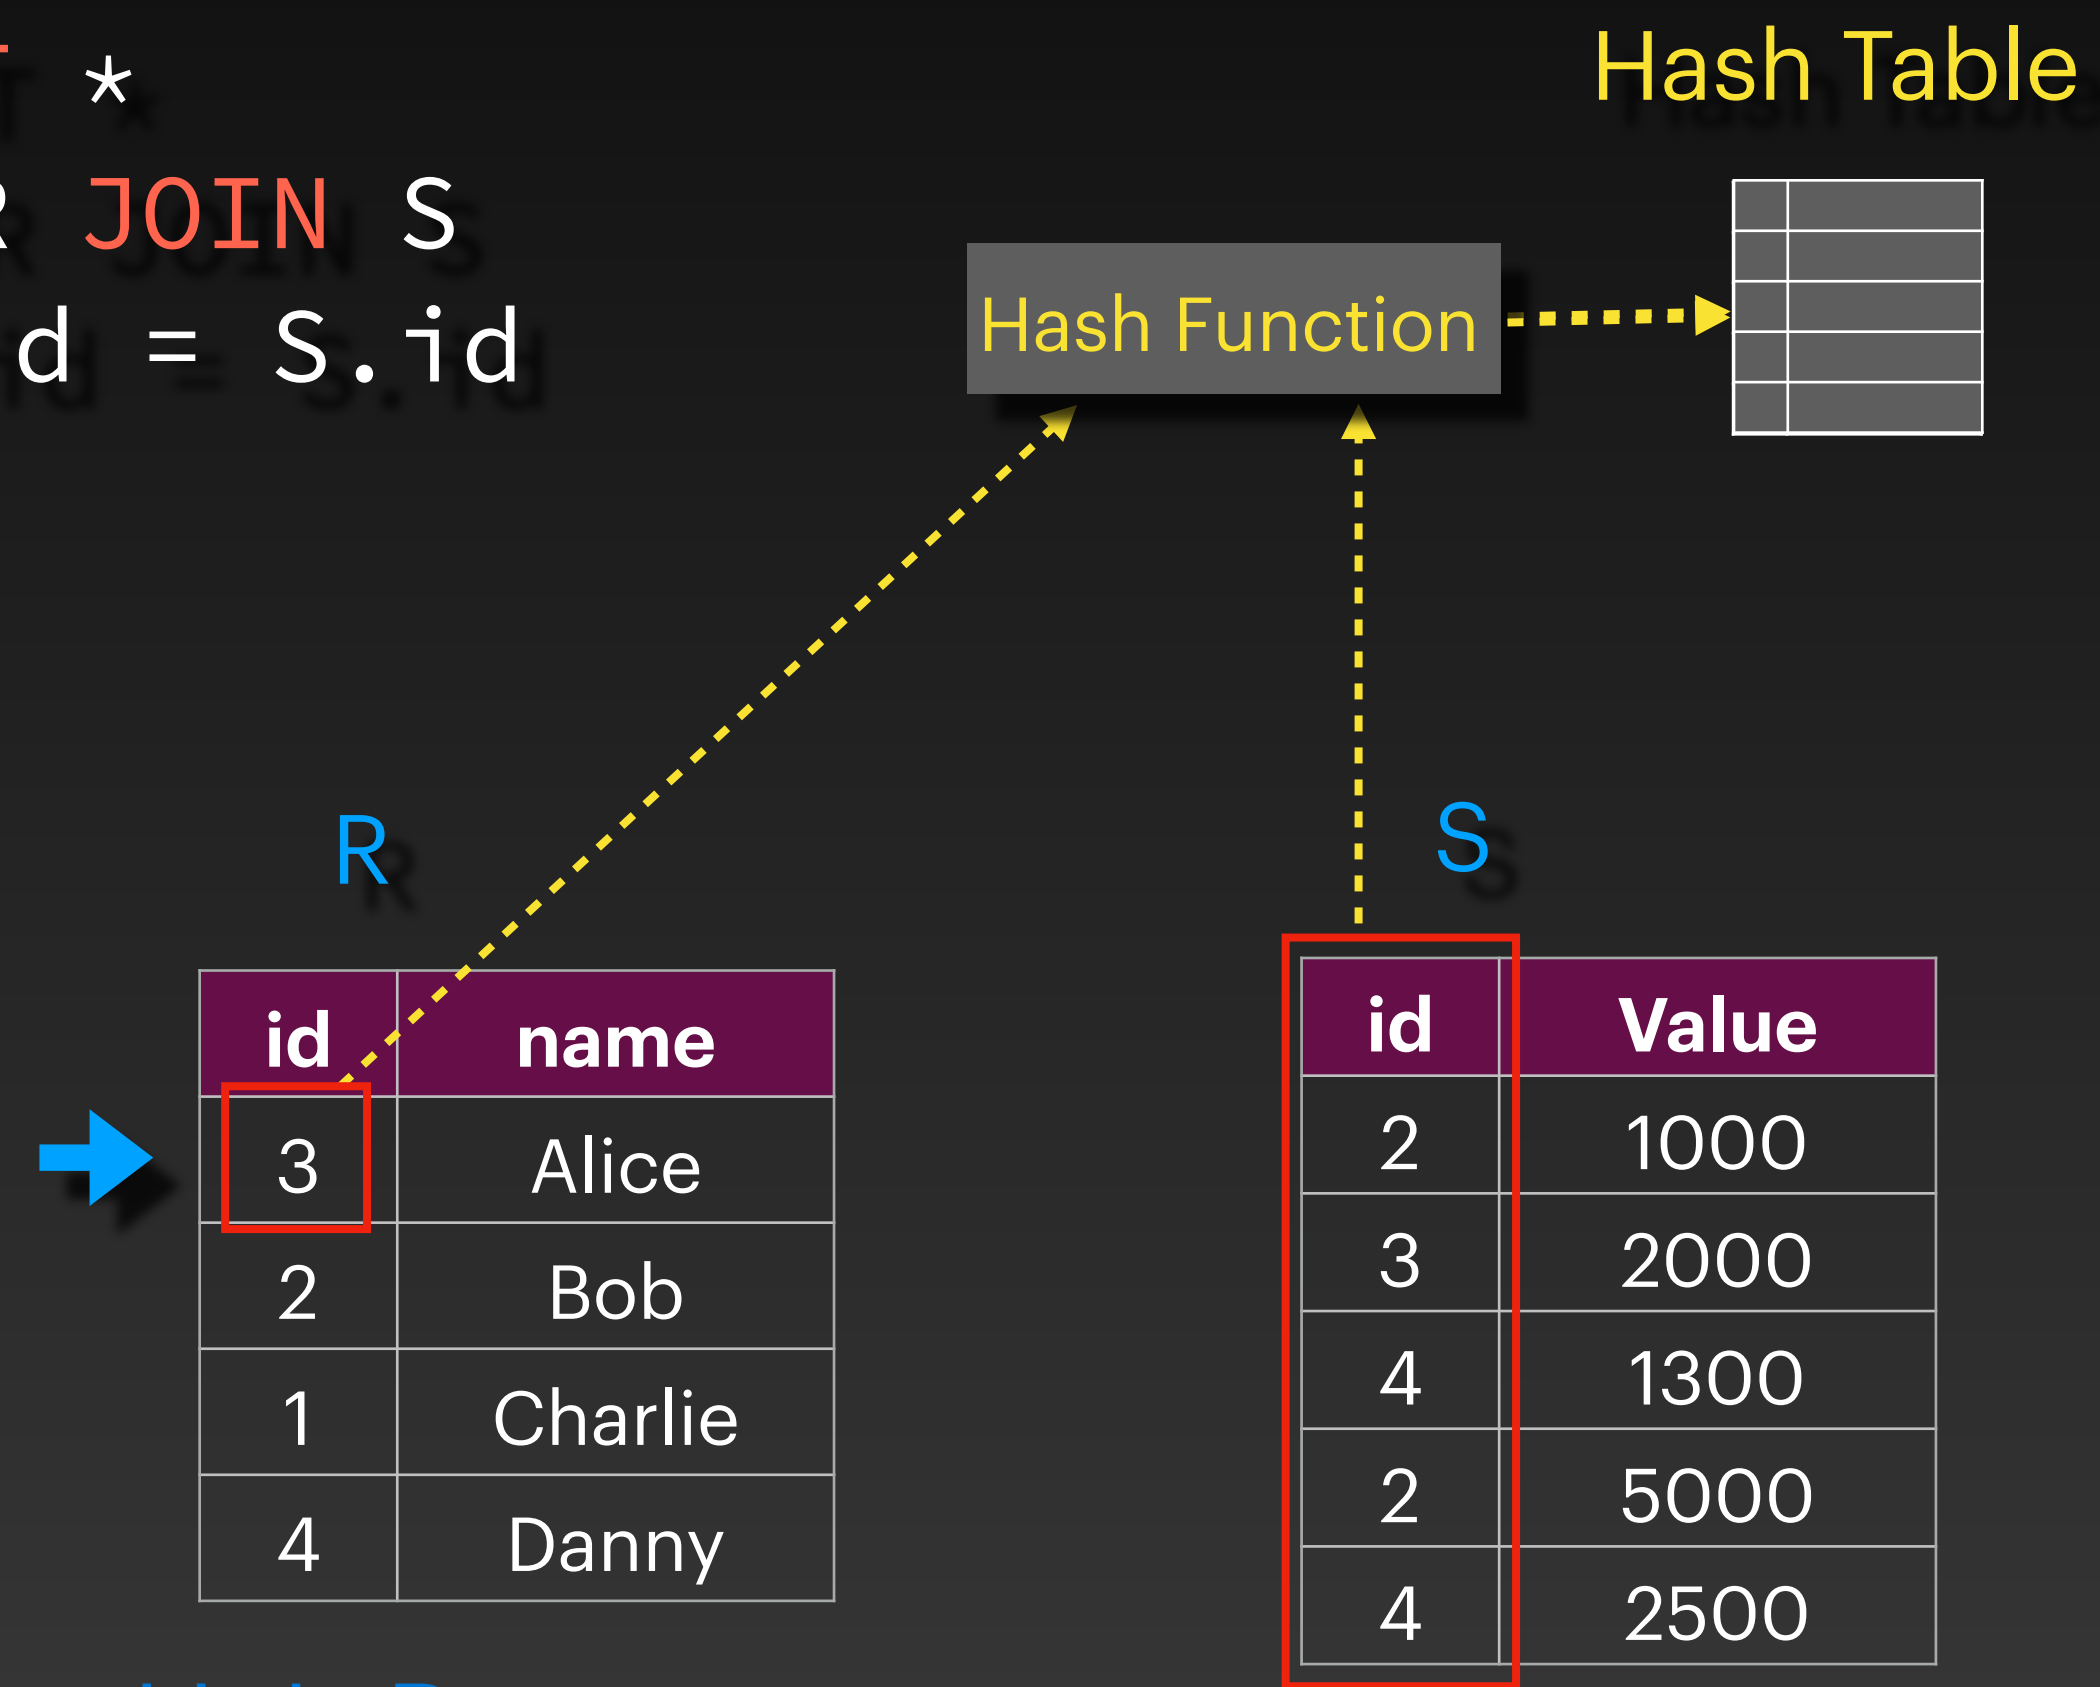
Cost = $(N_{\text{inner}} + N_{\text{outer}}) * \text{Cost_lookup_in_hash_table}$
consider outer, inner a big & small tables and optimizer
chooses the brst one to build a hash table for.

>> temp index dor this query inly.

Hash Join

Hash Join

```
SELECT *  
FROM R JOIN S  
ON R.id = S.id
```



iterate over ids in R.
hash the current id and
lookup its value in hash table

create hash table for ids in S

Output

R.id	R.name	S.id	S.value
3	Alice	3	2000
2	Bob	2	1000
2	Bob	2	5000
4	Danny	4	1300
4	Danny	4	2500