

SQL QUERY LIFE CYCLE

AMR ELHELW



CRUD Operations

C

- CREATE / INSERT

```
INSERT INTO employees (first_name, last_name, salary)
VALUES ('John', 'Doe', 60000);
```

R

- READ (SELECT)

```
SELECT first_name, last_name, salary
FROM employees
WHERE salary > 50000;
```

U

- UPDATE

```
UPDATE employees
SET salary = salary * 1.1
WHERE salary < 50000;
```

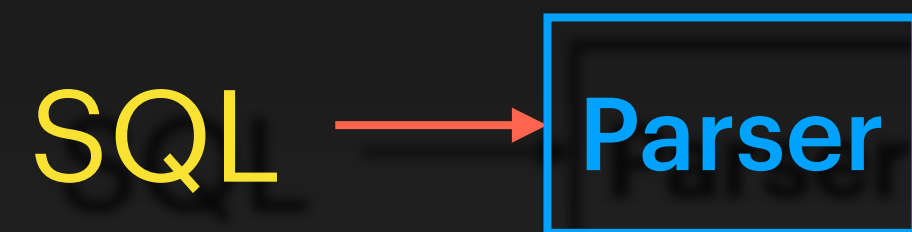
D

- DELETE

```
DELETE FROM employees
WHERE employee_id = 123;
```

Query Engine





```
SELECT id, name FROM user  
WHERE status = 'A' AND age > 20
```

is string for system and the engine converts words
so tokens stored in the system to be able to work with them
every unknown word would be considered an identifier
even if it was wrong
in this phase no checks for identifiers. it happens in analyzer

Parsing



SELECT name, phone FROM user WHERE status = 'active' AND agee > 20

SELECT

Ident
name

,

Ident
phone

FROM

Ident
user

WHERE

Ident
status

=

Const
'active'

AND

Ident
agee

>

Const
20

1
Tokens

Keywords	SELECT - FROM - AND - OR - WHERE - NULL - NOT - IN - GROUP - ...
Operations	> - < - = - <= - >= - + - ...
Other symbols) - (- , (comma) - . (dot) - ...
Constant values	'Tom' - 16 - '02-04-2023' - ...
Identifiers	Anything else

2 Syntax Checking

```
SELECT name, phone FROM user WHERE status = \active AND agee > 20
```



Missing '

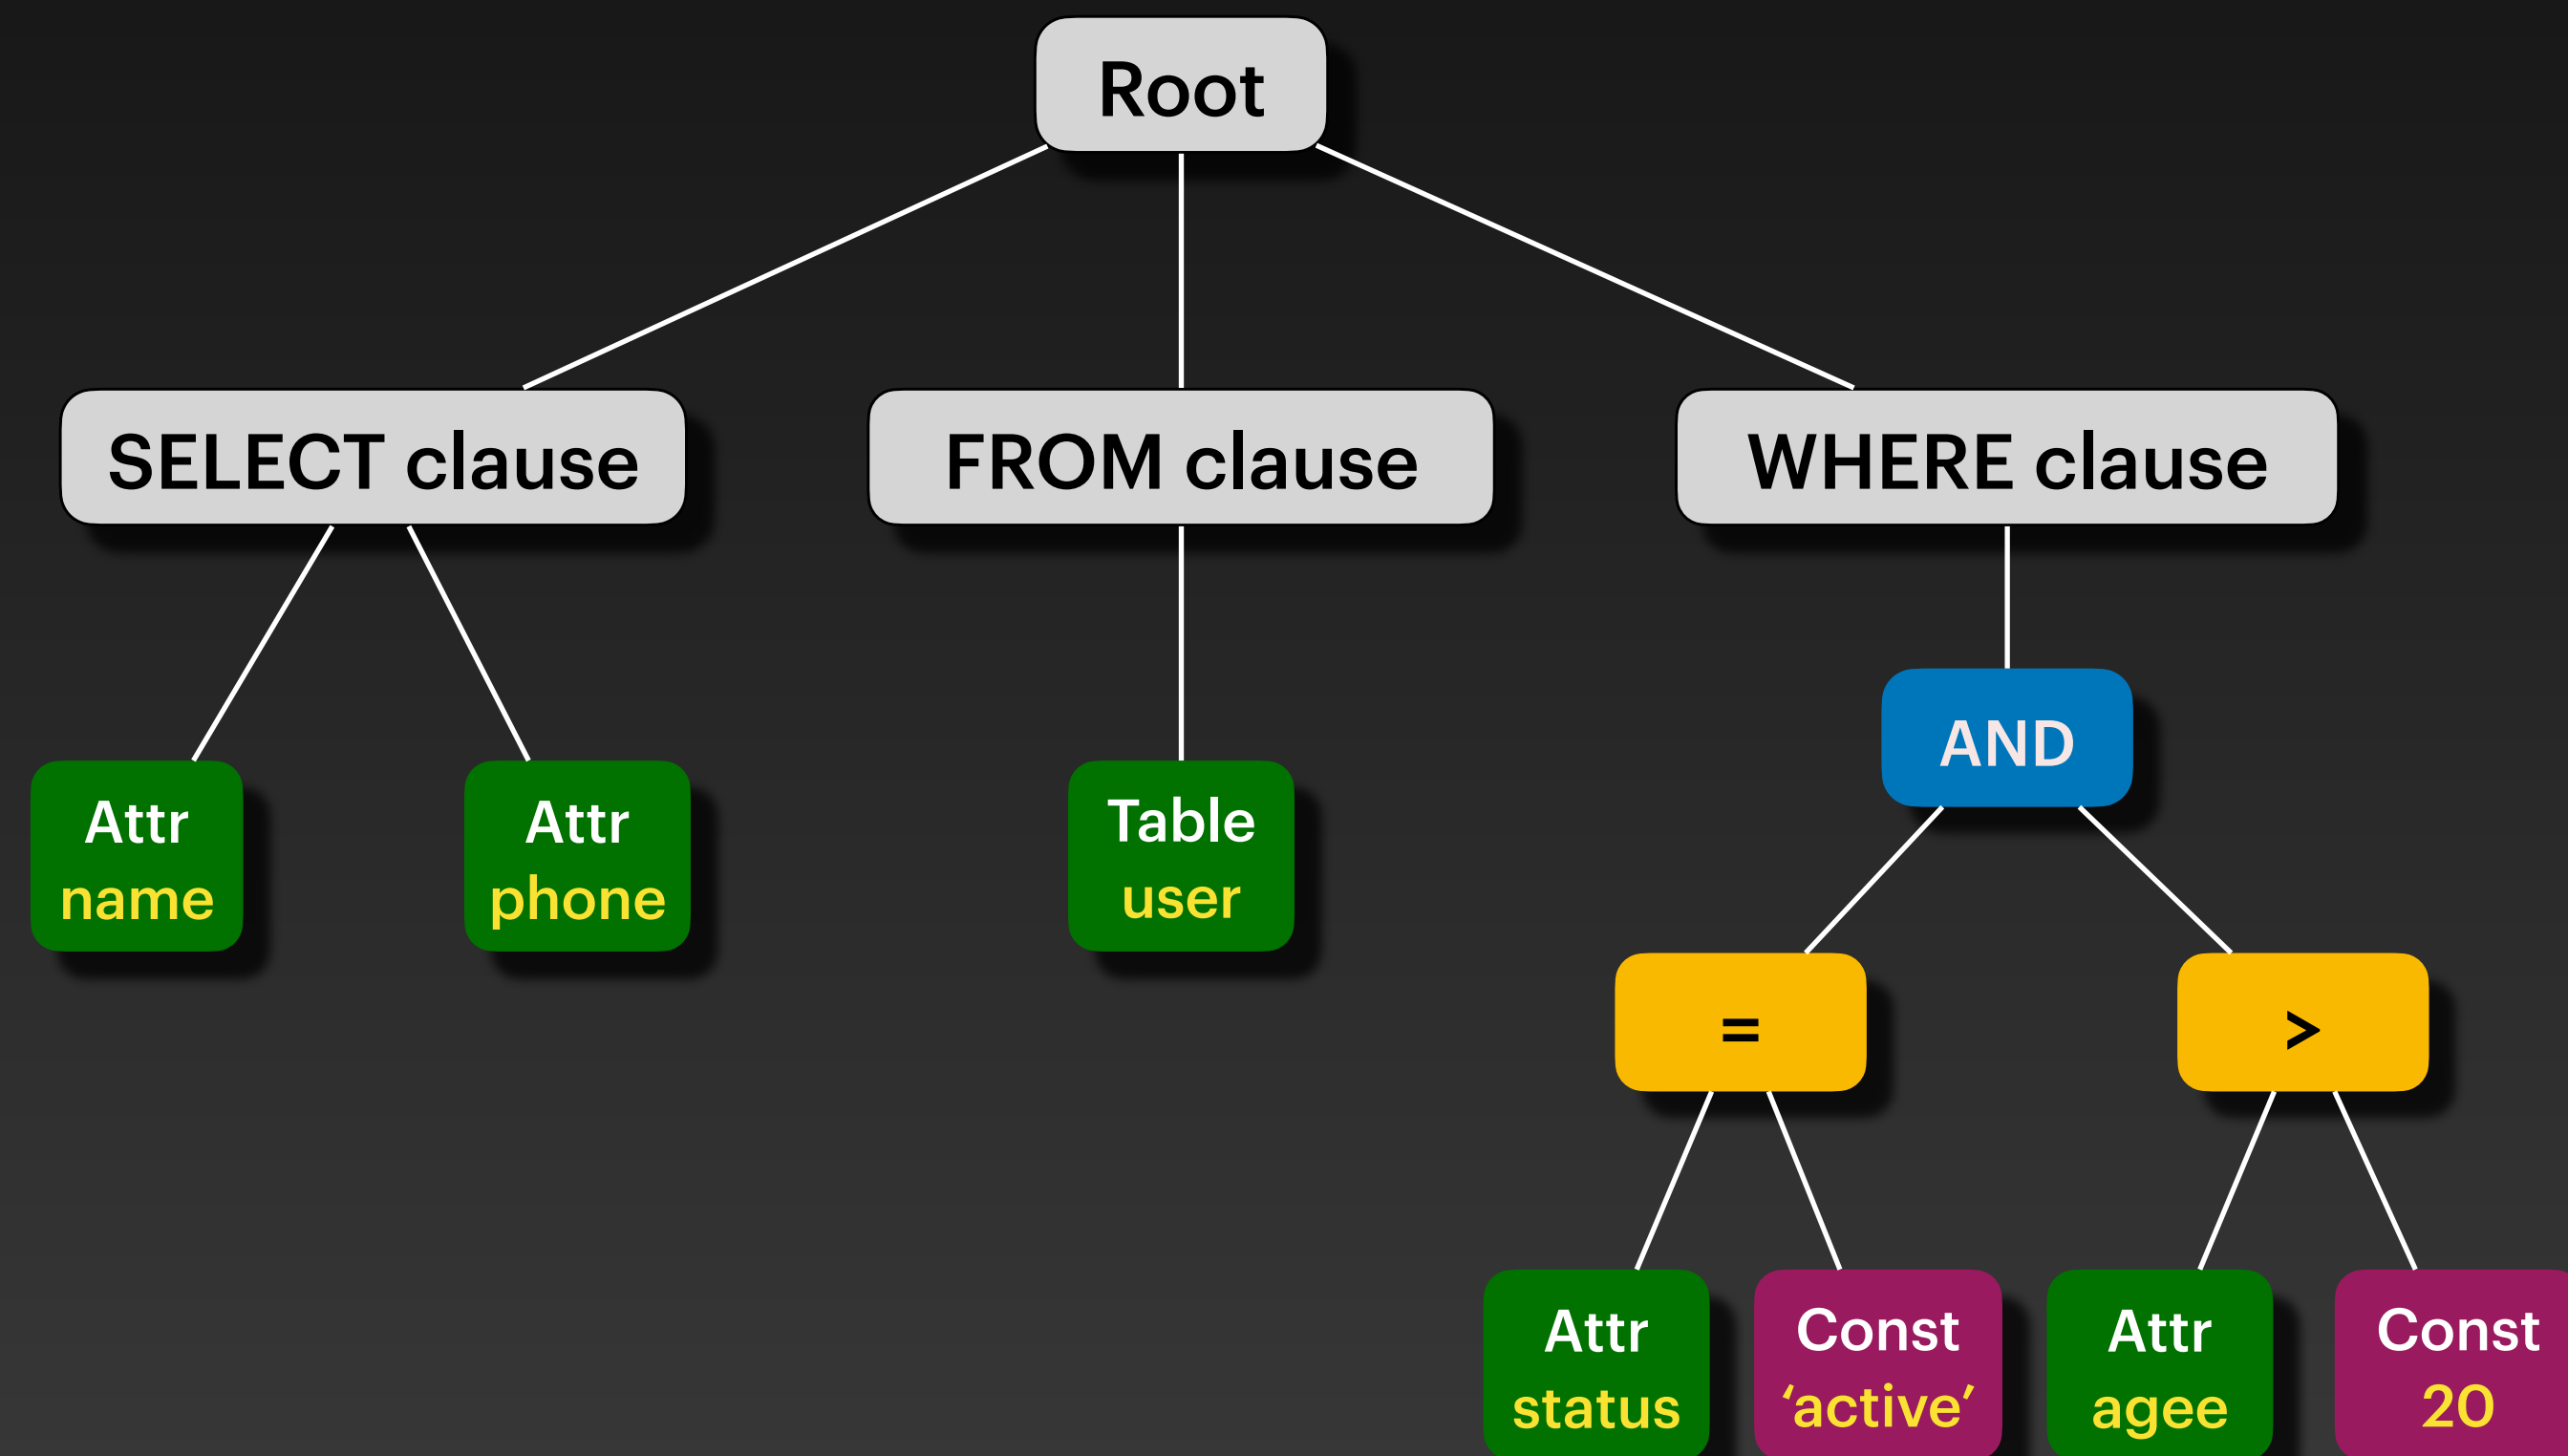
```
SELECT name, phone WHERE status = 'active' AND agee > 20
```



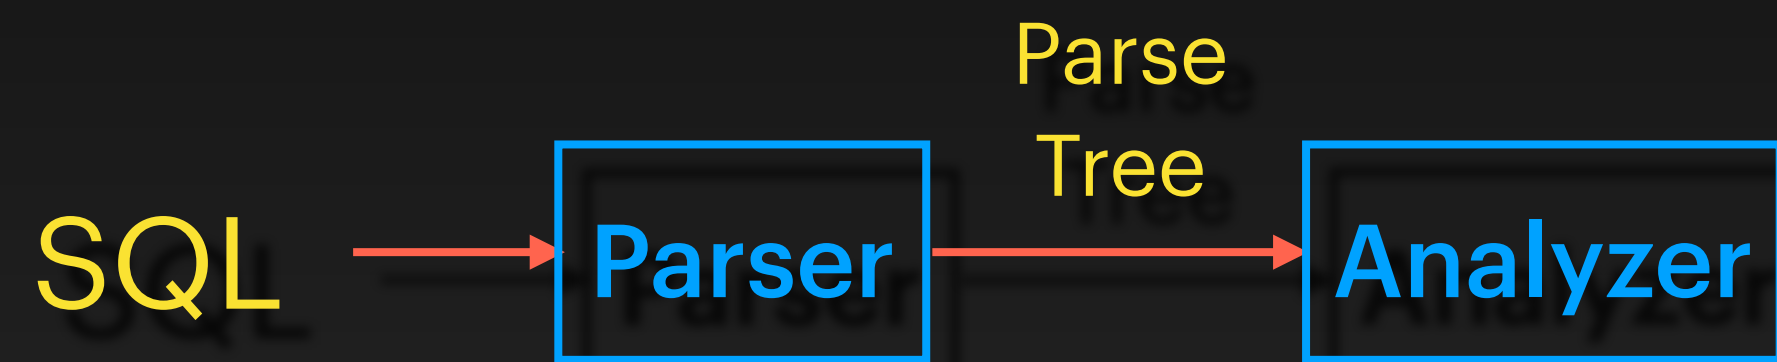
Missing FROM

3 Parse Tree

SELECT name, phone FROM user WHERE status = 'active' AND agee > 20

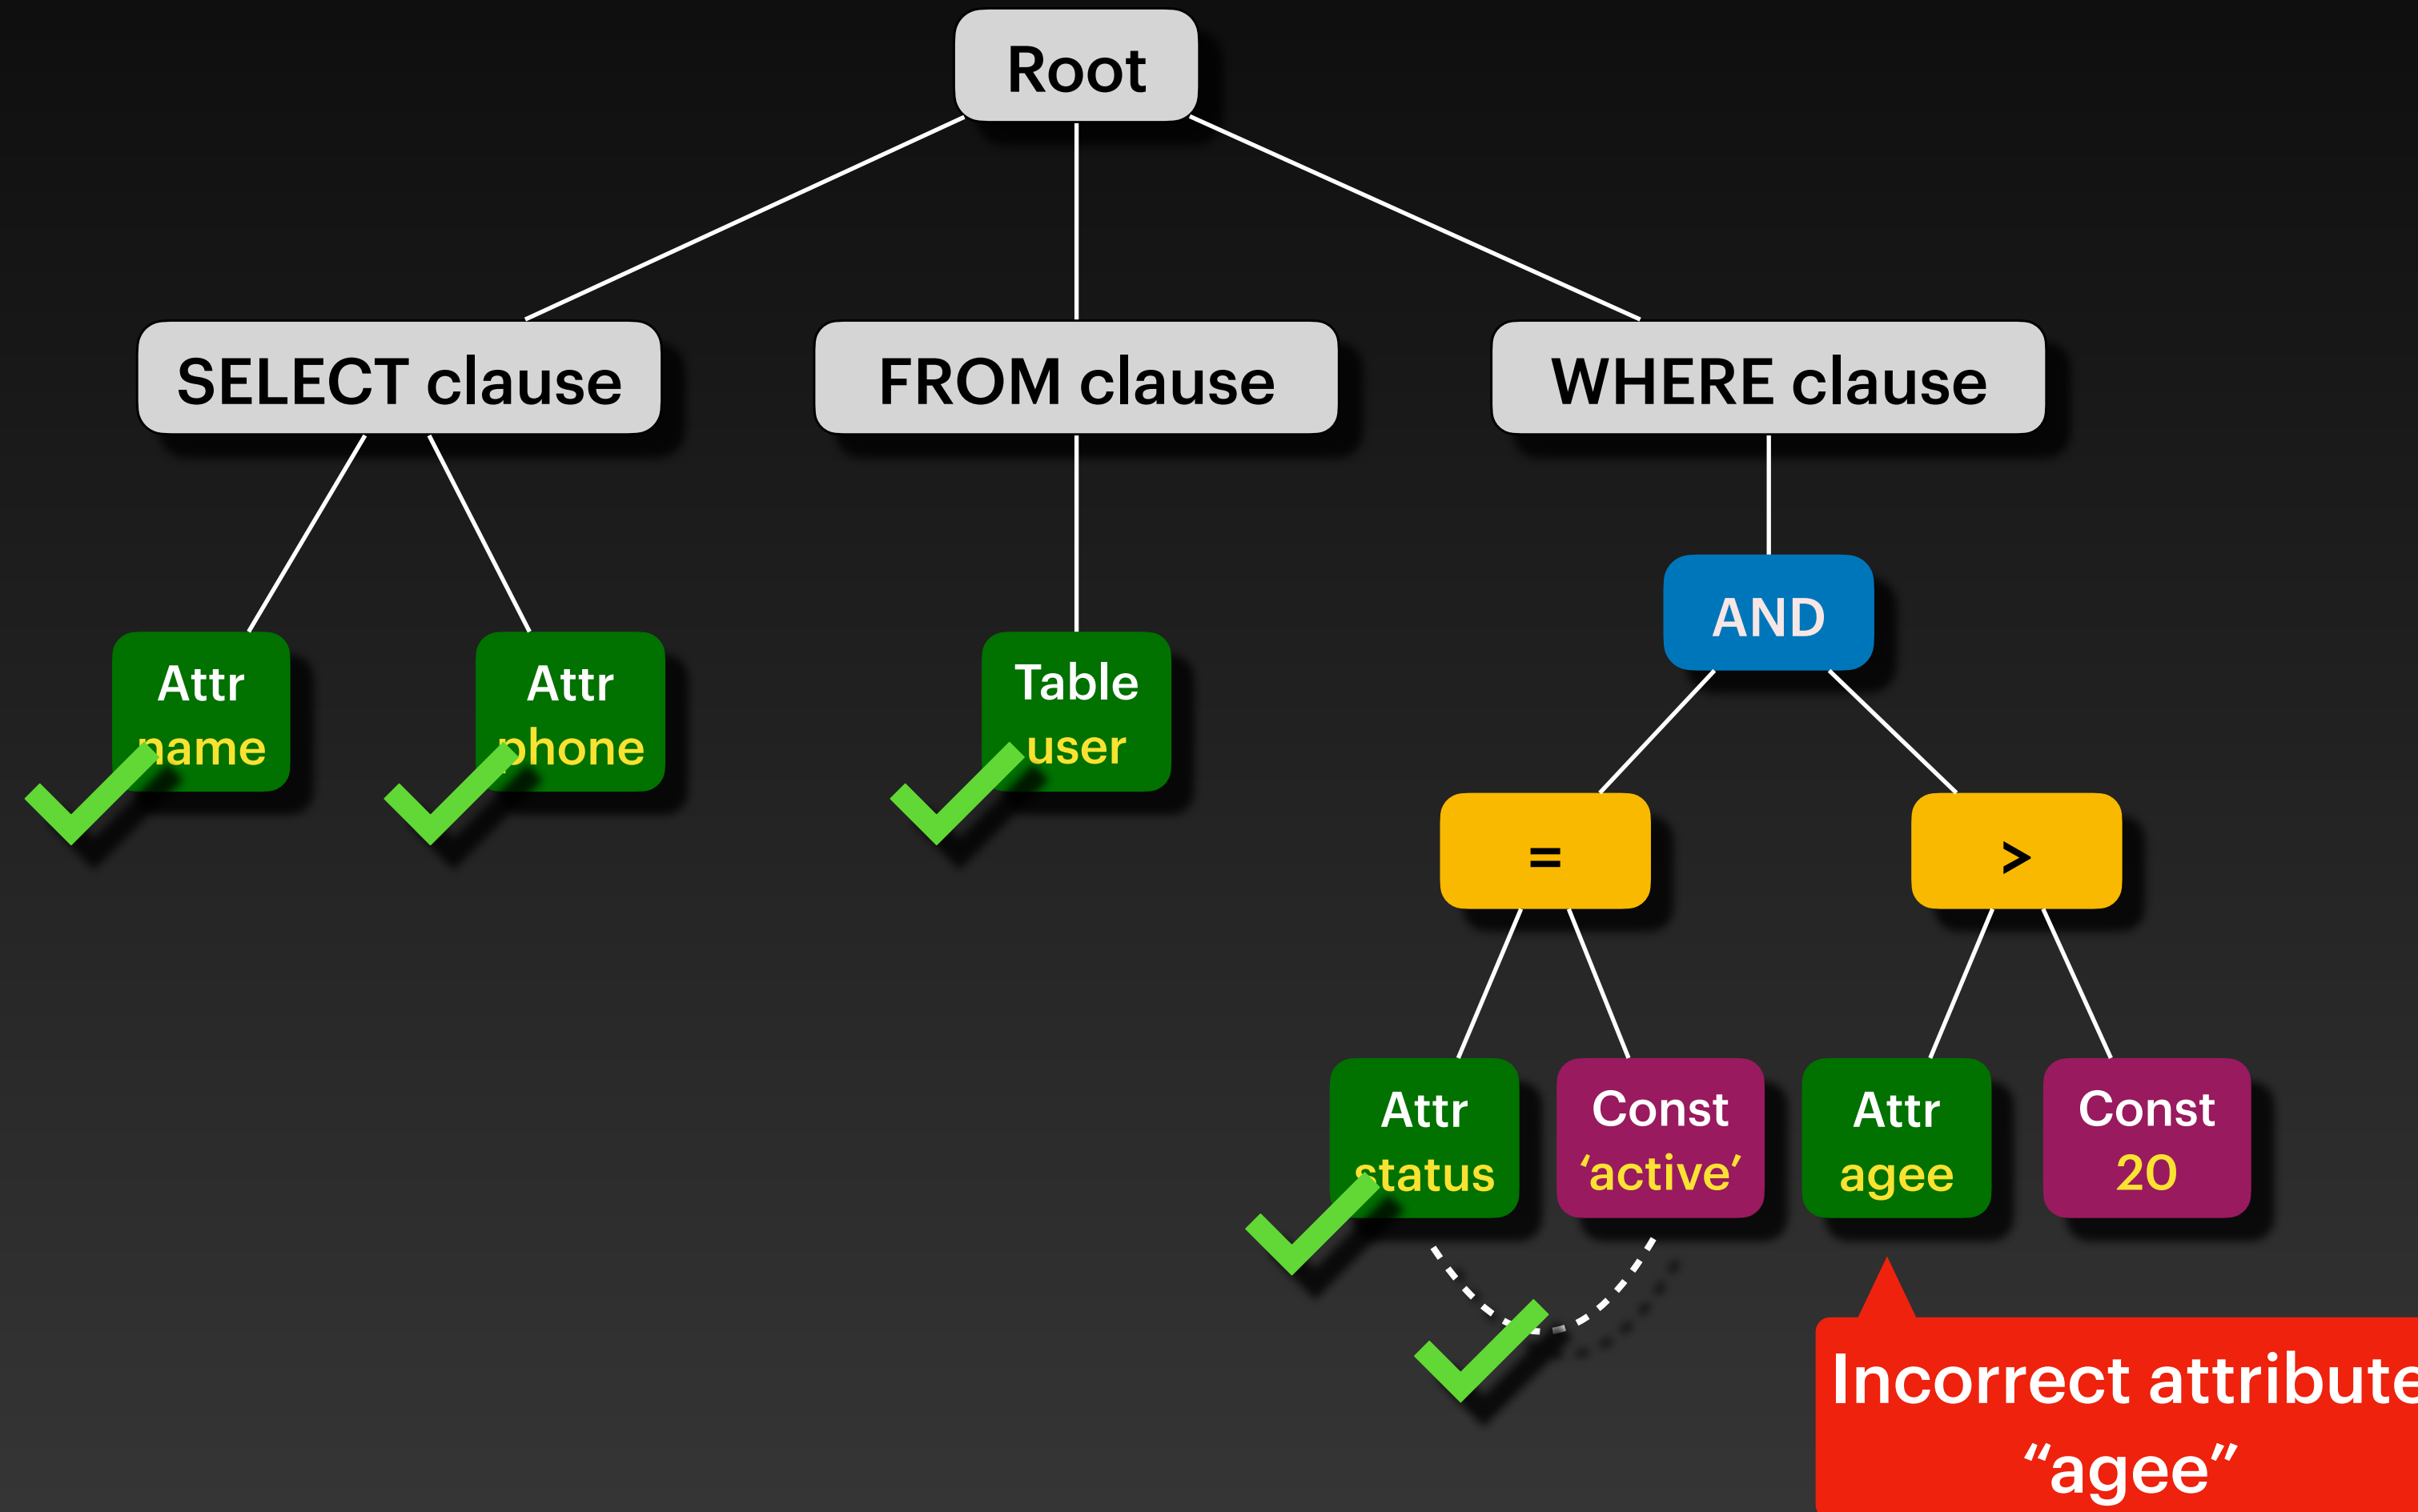


Query Engine



Query Analysis

Parse Tree



analyzer steps:

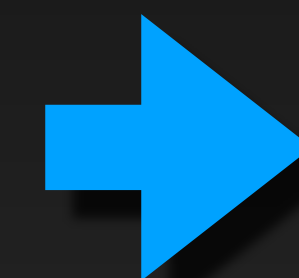
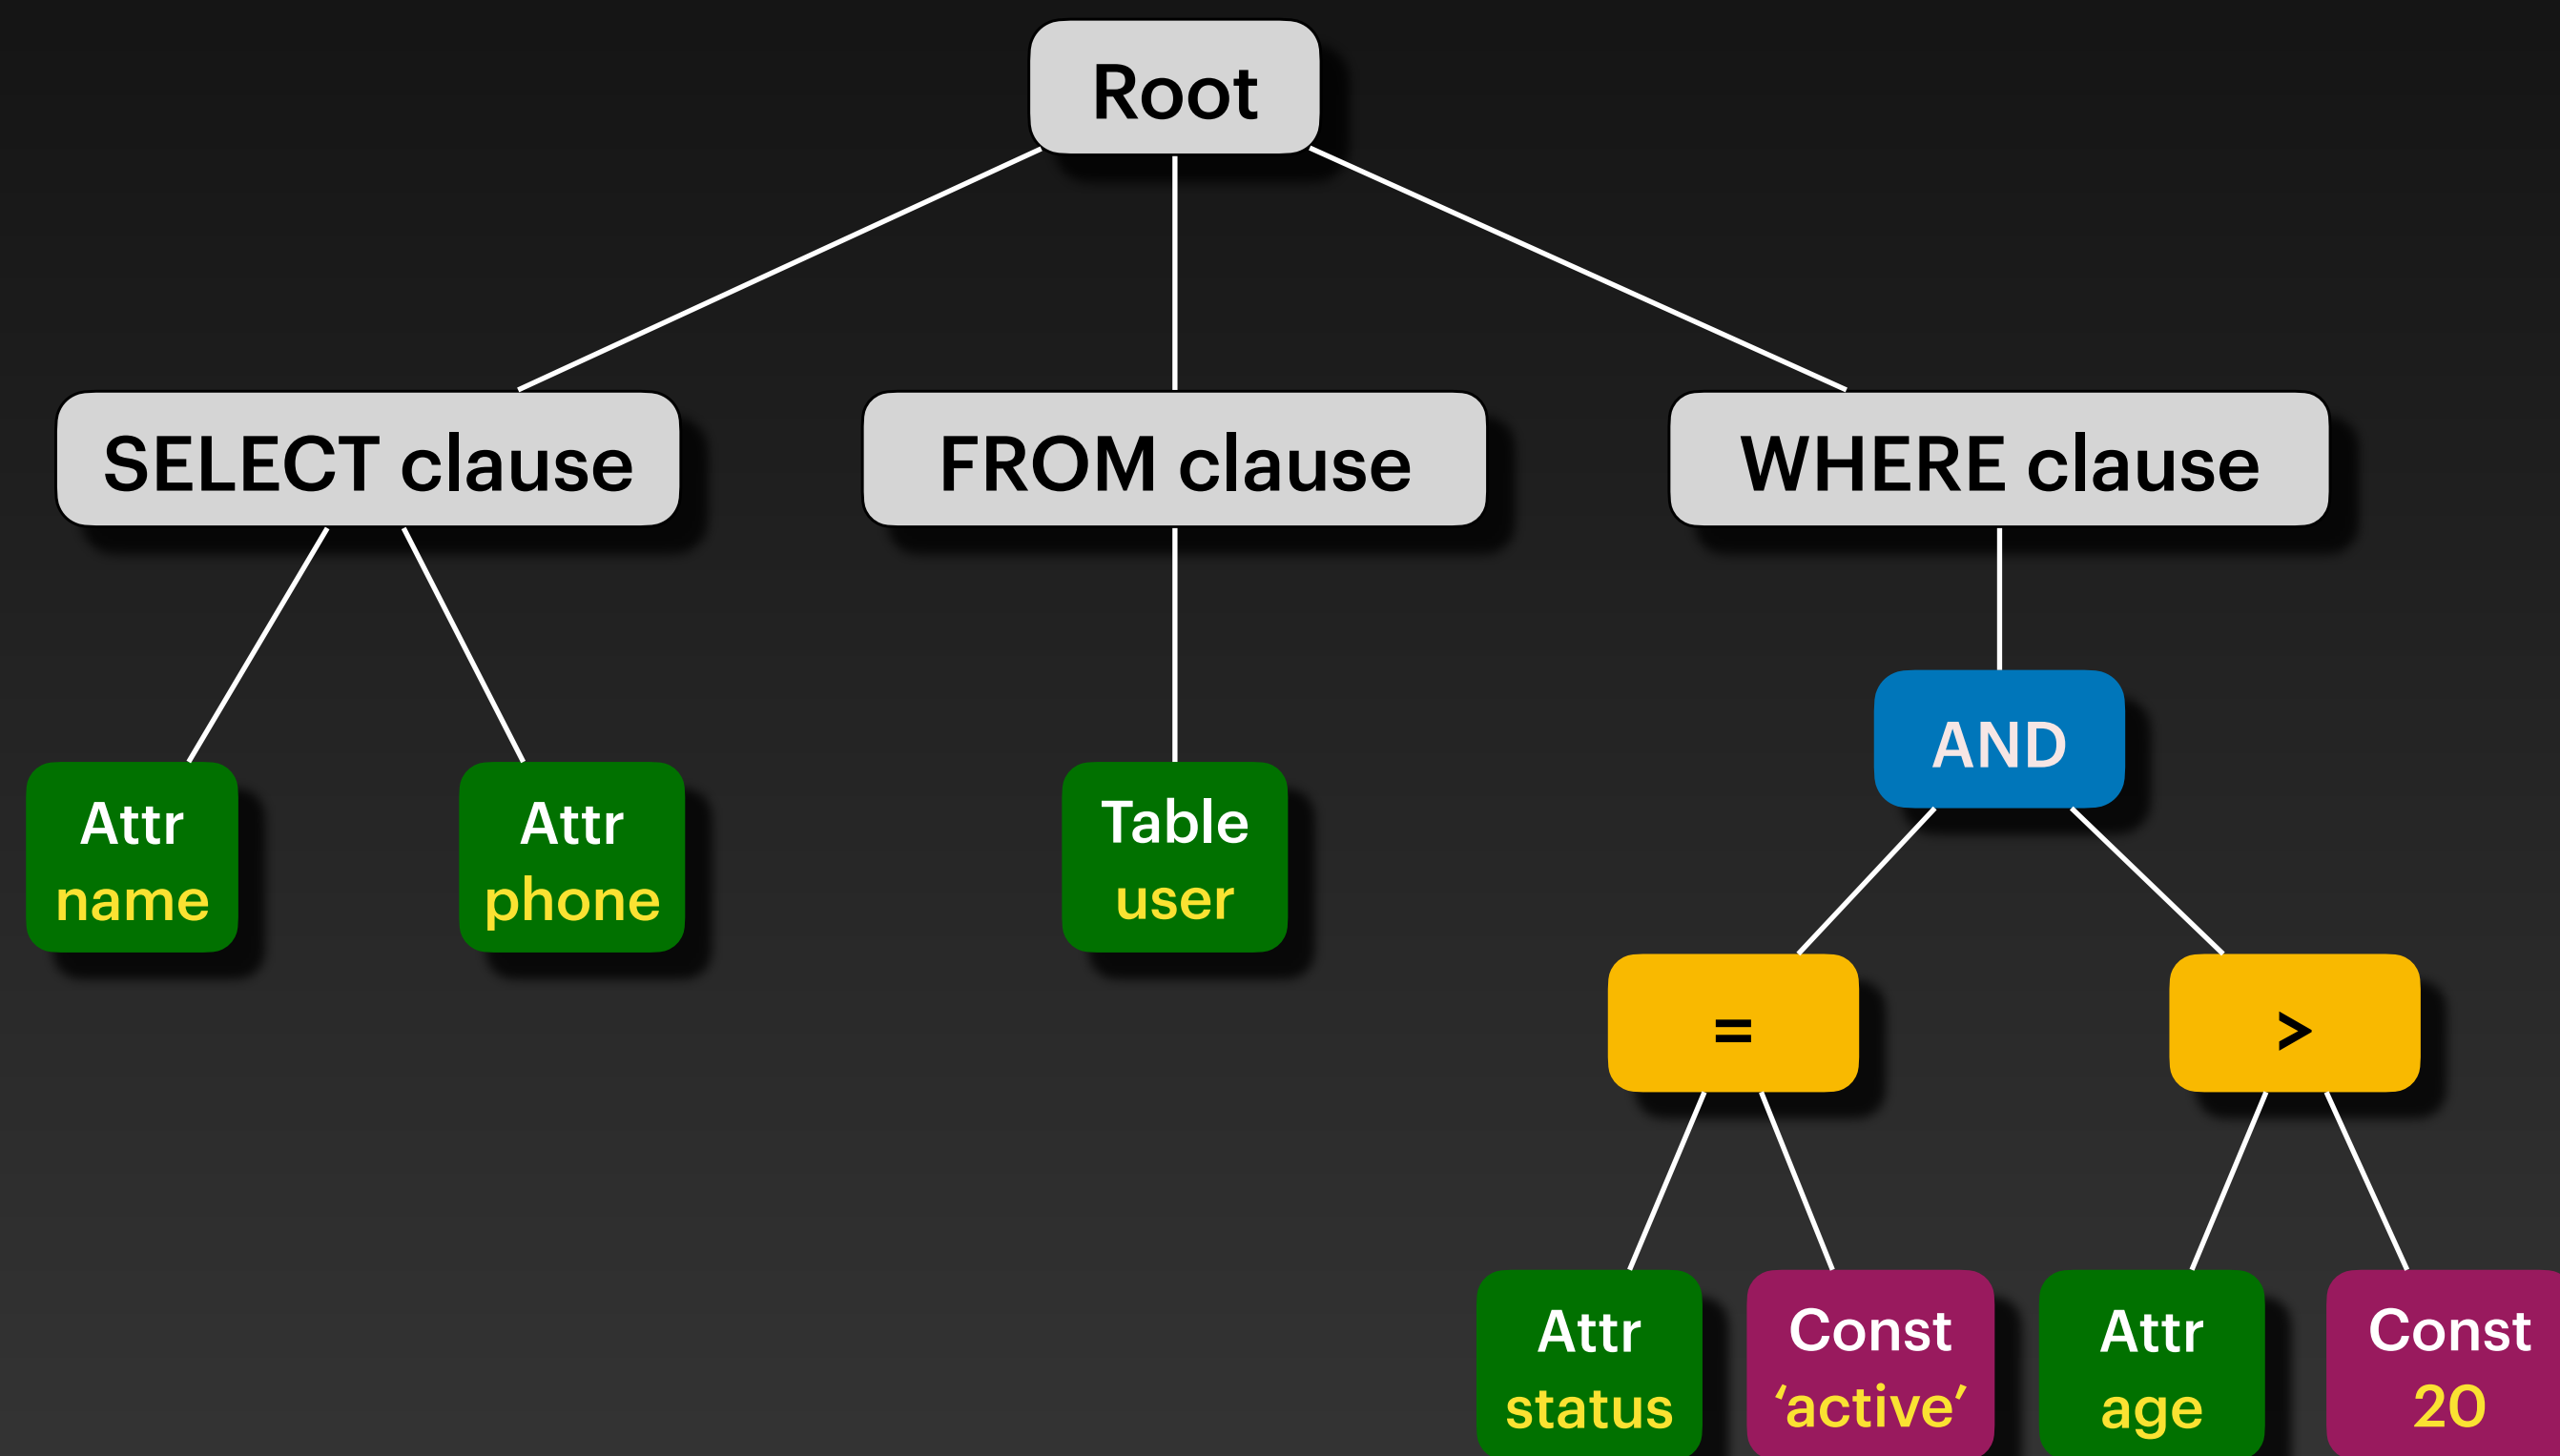
1. semantic checks:

validates tables, cols, data types in comparison, ambiguous names, functions

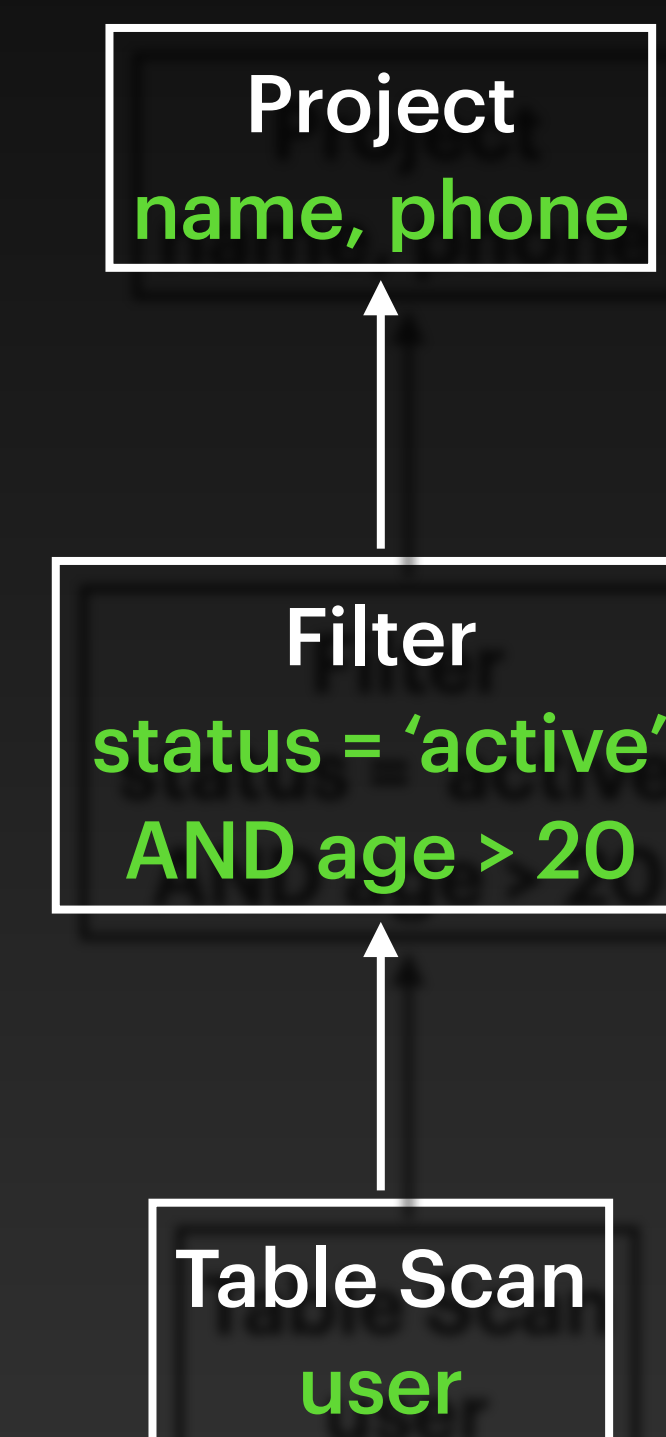
2. Initial query plan (not optimized, such as linear algebra)

Query Analysis

Parse Tree

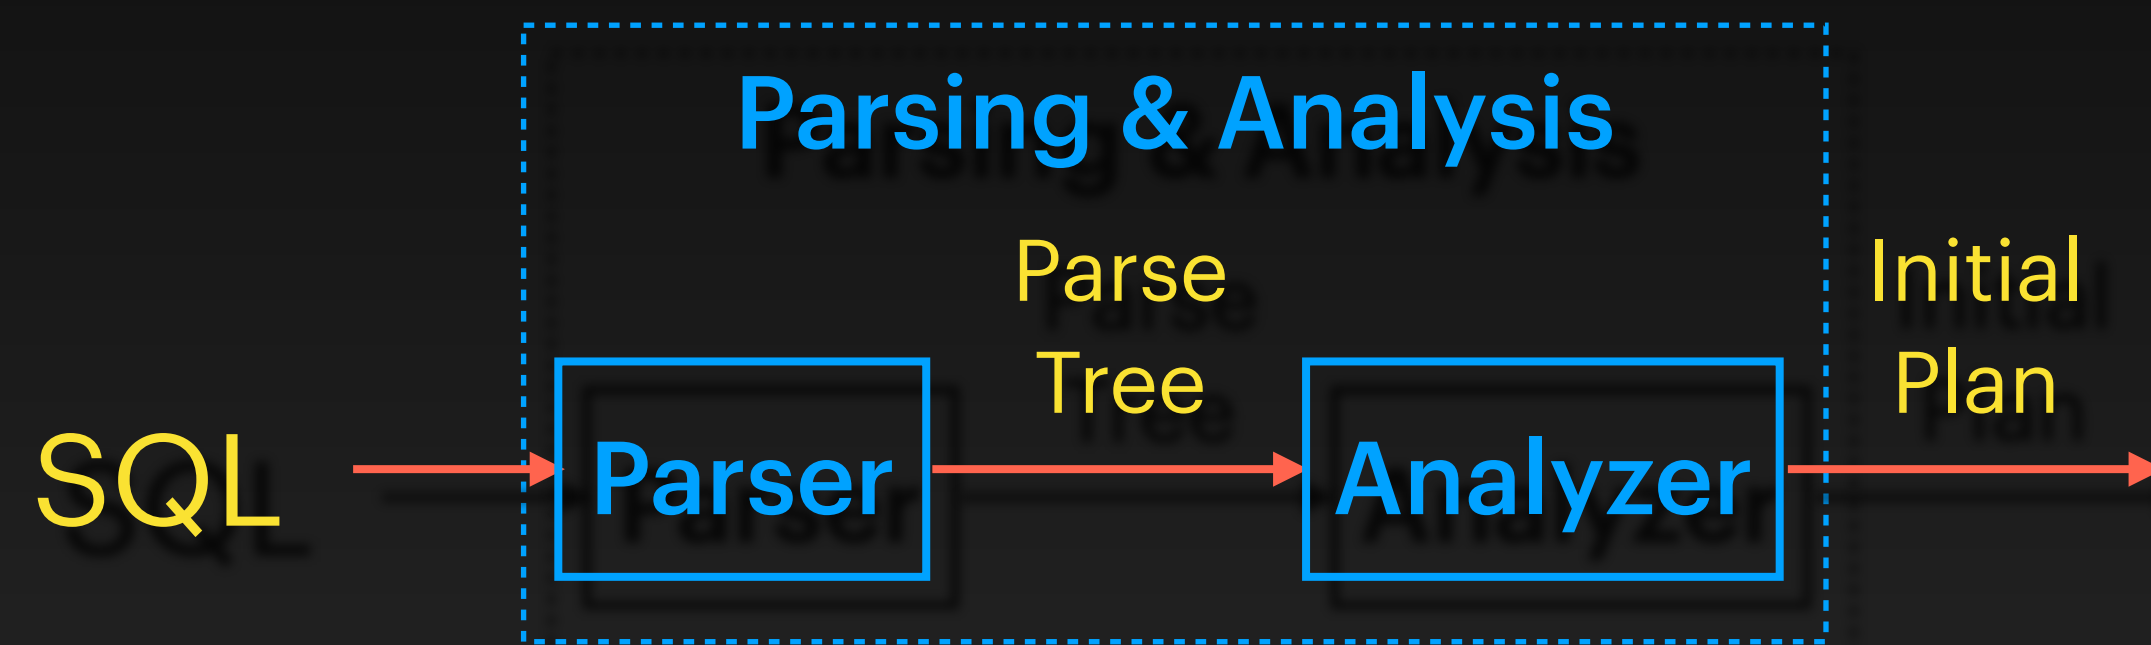


Initial Query Plan



Query Engine

Catalog: is a namespace collection of metadata about tables, attributes, indexes and so on. they are tables also.



pg_class table: tables metadata

pg_attribute table: attributes/cols metadata

(has attrelid is foreignkey from pg_class table)

pg_type table: data types metadata

pg_namespace: namespaces metadata

(such as public(pg users), catalogs in pg_catalog)

pg_stats table: statistics table contains data such as:
num of rows, distinct, nulls, avg size...etc

explain: gets the query plan

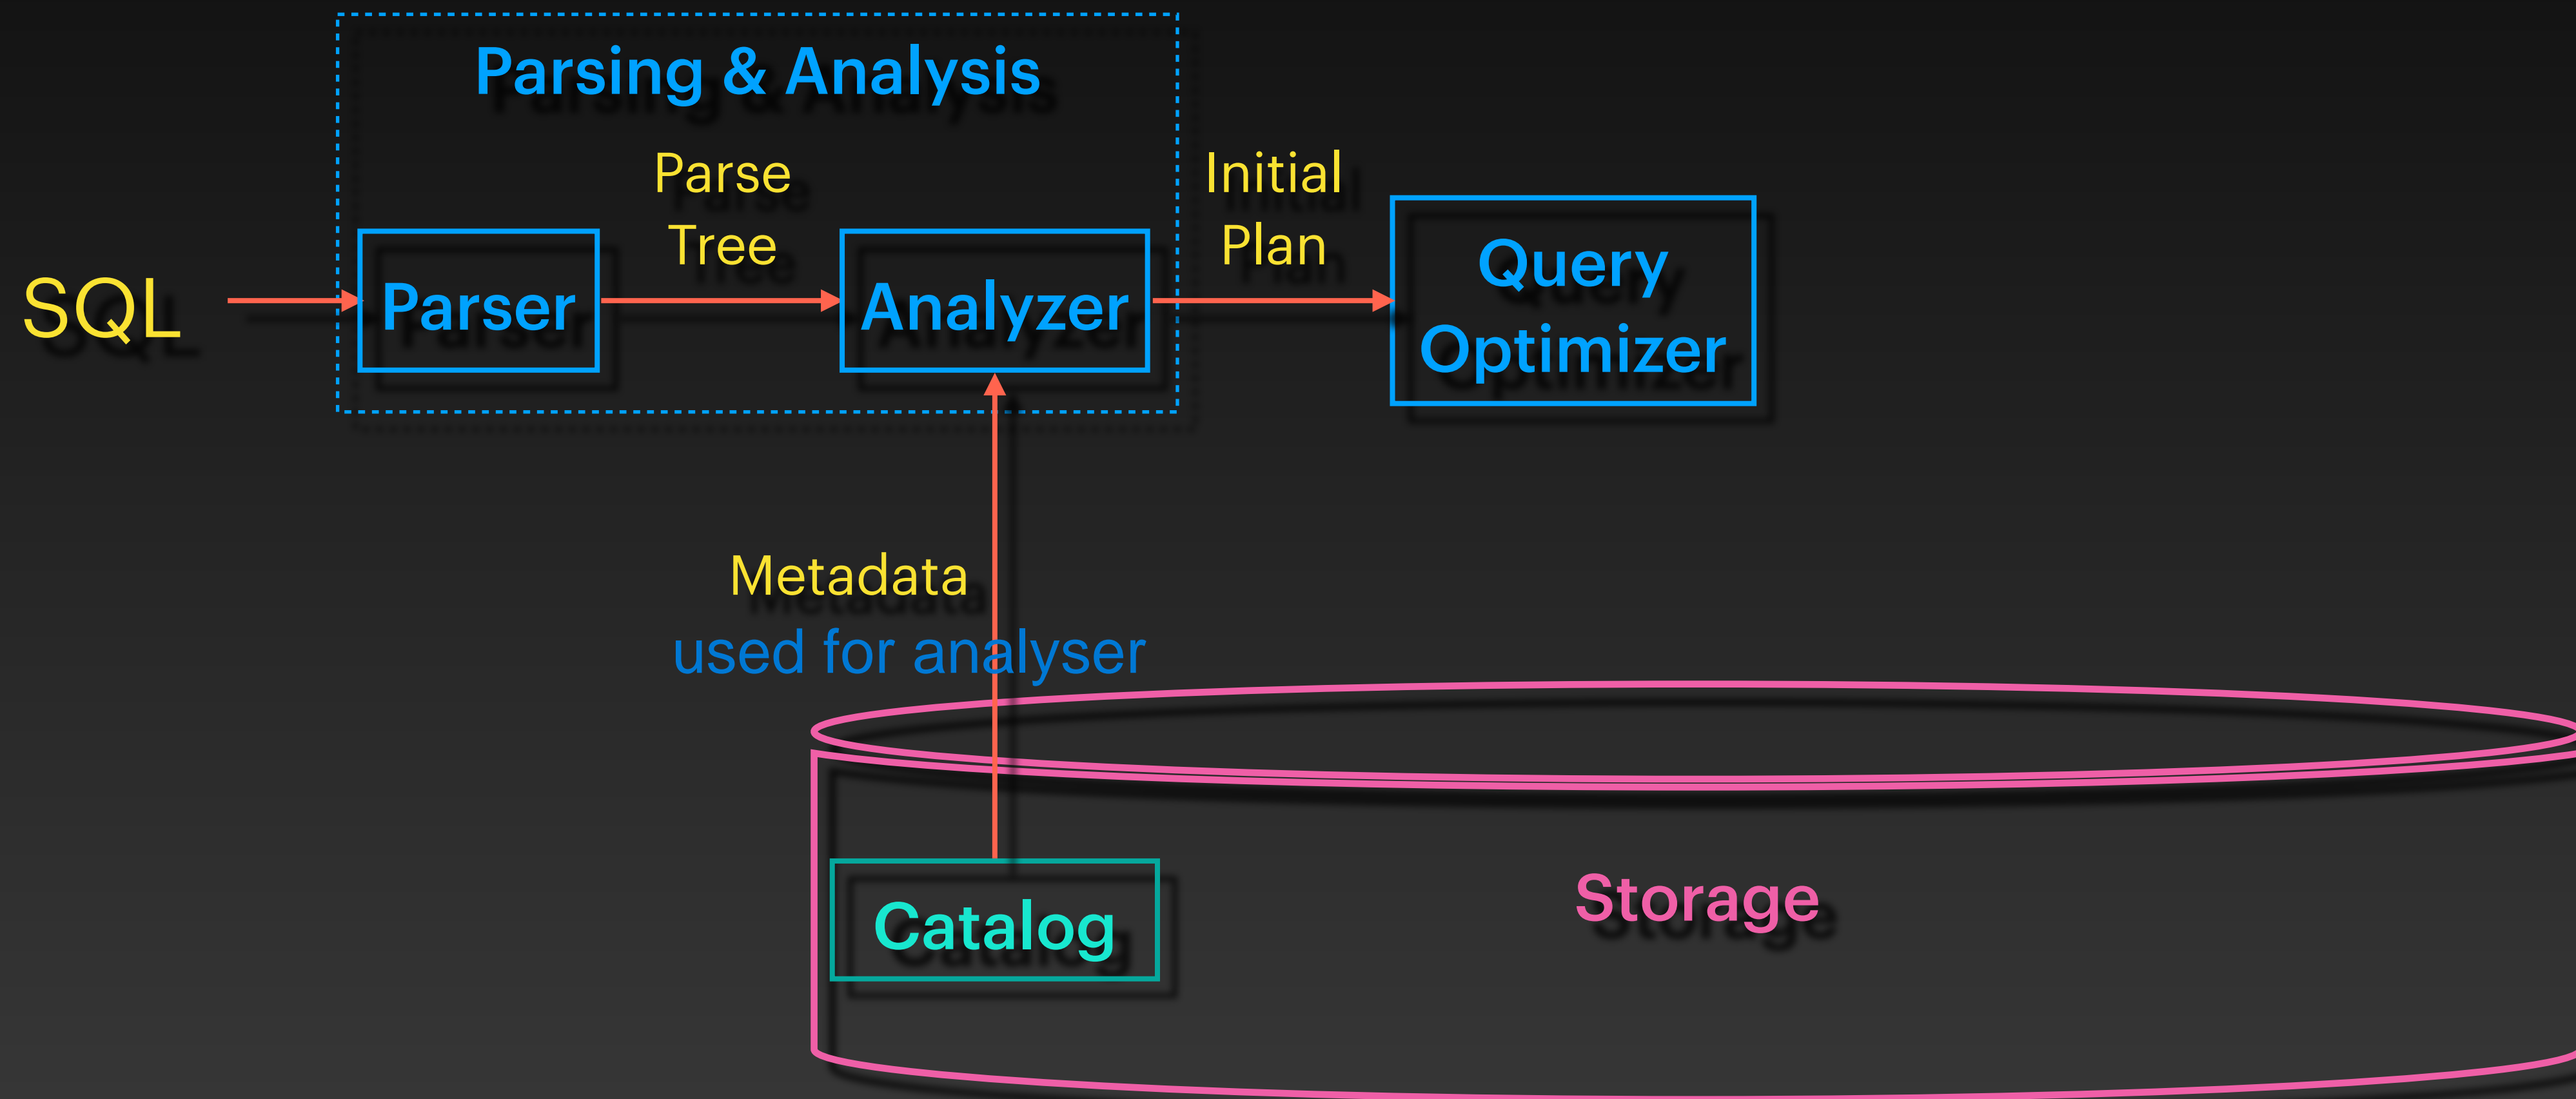
explain verbose: gets the returned attributes from each step in the plan

explain analyze: gets estimated and actual time and num of rows

tips:

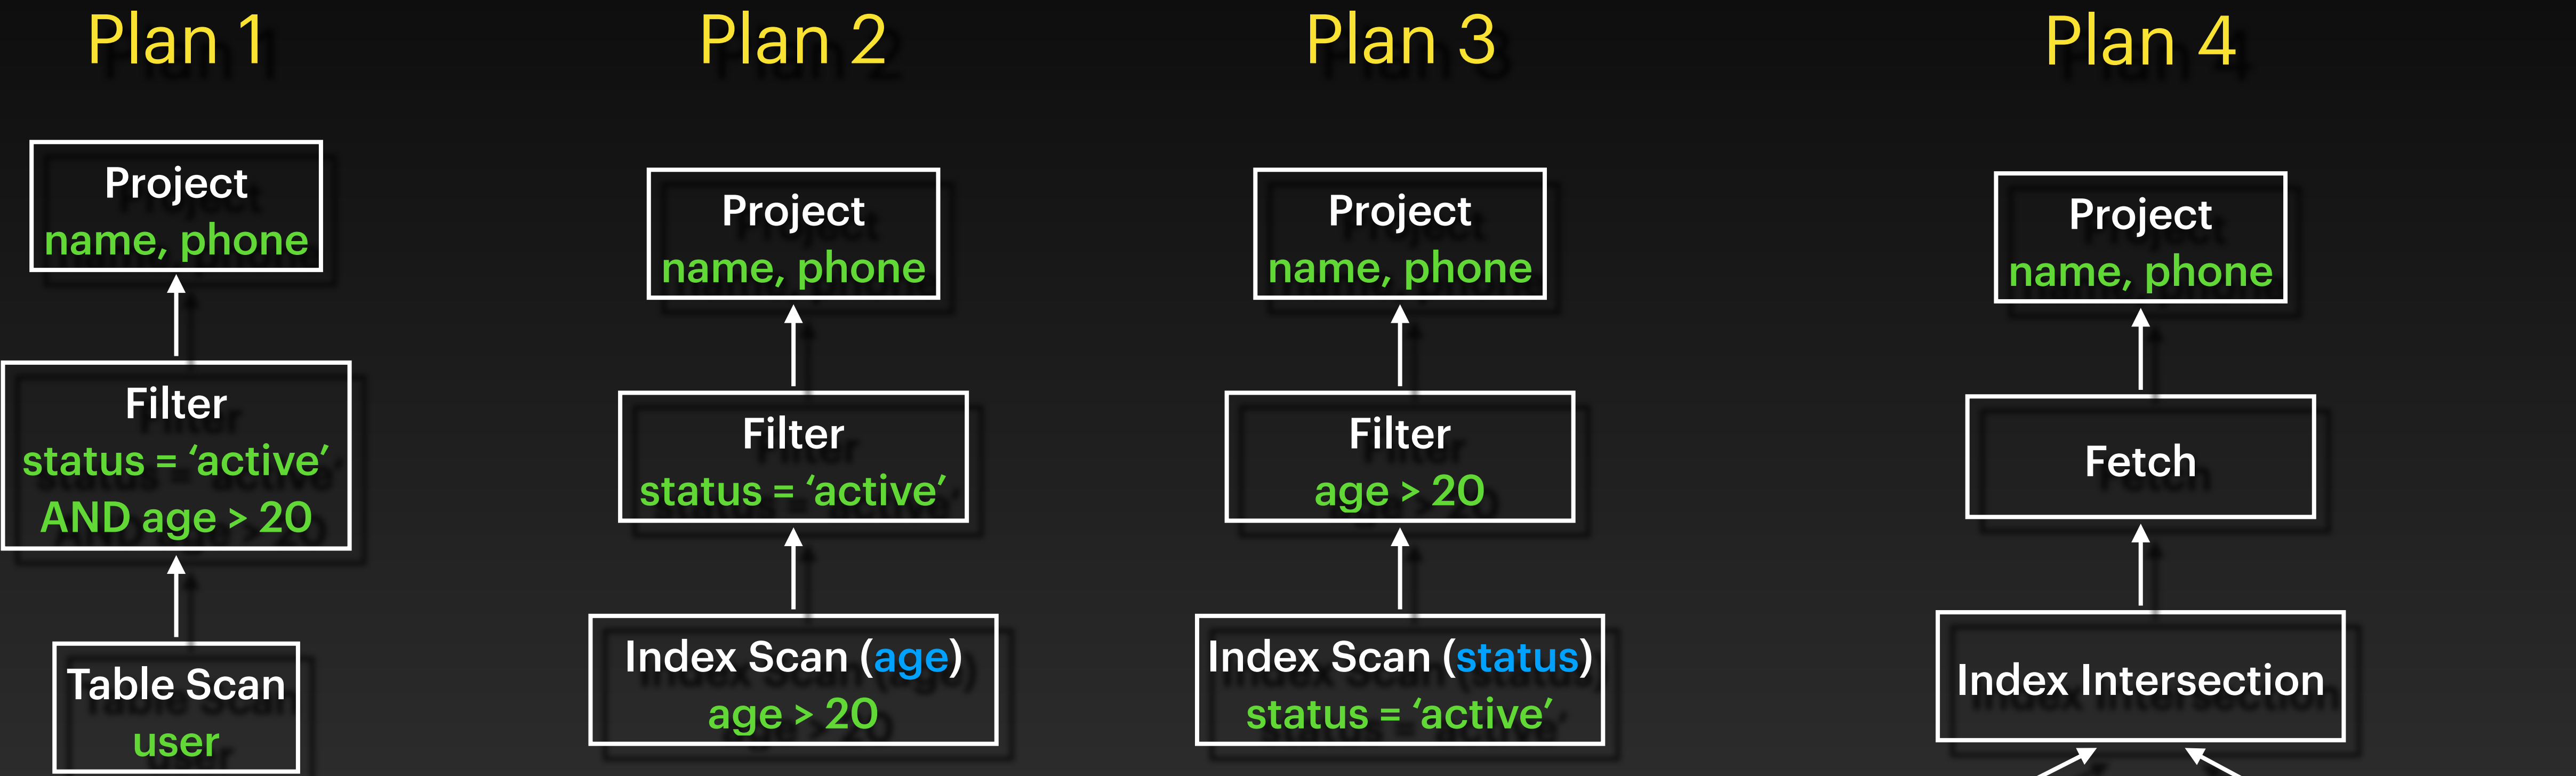
- * pgexplain website visualize the query plan from explain.
- * plan are read from bottom to top
- * you can format the output of explain as xml or json.

Query Engine



there is a limitless number of combinations of query plans
that no optimizer can calculates all of them which makes it not totally efficient.

Query Optimizer



Indexes:

Index on (age)

Index on (status)

gets alternative equivalent plans according to the provided metadata in catalog.
then compares them using the cost model which is formulas calculations for each type of queries.
database statistics: metadata tables contains statistics about each table, number of rows..etc

Query Engine

